

# HorseIR : An Array-based Approach to SQL Queries


Hongji Chen

August 30, 2017

# What is HorseIR?



- ▶ array-based programming language
- ▶ uniform intermediate representation(IR)
- ▶ optimization framework



# What is HorseIR?



# In-memory Database System and IR-based Query Engine



- ▶ row-oriented or column-oriented systems
  - ▶ row-oriented storage: run as few operation as possible on row record (most database use-case, e.g. logging)
  - ▶ column-oriented storage: good for set operations for whole table data (e.g. analytics)
- ▶ complex primitives set or reduced primitives set
- ▶ user-defined functions (UDF)
  - ▶ not in SQL standard, but as a language extension in most systems
  - ▶ flexibility?
  - ▶ optimize potential?
- ▶ parallel code generation

# Related Work - In-memory Database Systems



- ▶ SQLite
  - ▶ <https://sqlite.org/opcode.html>
- ▶ MonetDB
  - ▶ Stratos Idreos, Fabian Groffen, Niels Nes, Stefan Manegold, K. Sjoerd Mullender, and Martin L. Kersten. 2012. MonetDB: Two Decades of Research in Column-oriented Database Architectures. IEEE Data Eng. Bull. 35, 1 (2012), 40–45.
  - ▶ <http://sites.computer.org/debull/A12mar/monetdb.pdf>
- ▶ KDB+ System
  - ▶ <https://kx.com>



## ► SQLite Bytecode

- scalar based IR design
- dynamically typed
- no optimization on bytecode level
- designed for register based virtual machine



- ▶ pioneer in IR based execution engine design
- ▶ MonetDB Assembly language(MAL)
  - ▶ scalar based IR design
  - ▶ statically strongly typed, without sub-typing
  - ▶ encapsulate atomic relational algebra operation in each instruction
  - ▶ frontend translate queries into MAL
  - ▶ backend interpret MAL to generate result
  - ▶ optimizer optimize MAL to MAL



- ▶ K and Q programming language
  - ▶ fully functional programming (not IR-based)
  - ▶ array-based language design (array/vector are first-class type)
  - ▶ focus on data analytics
  - ▶ powerful and efficient primitives implementation
  - ▶ limited intra-procedural or inter-procedural optimization



# Related Work Summary

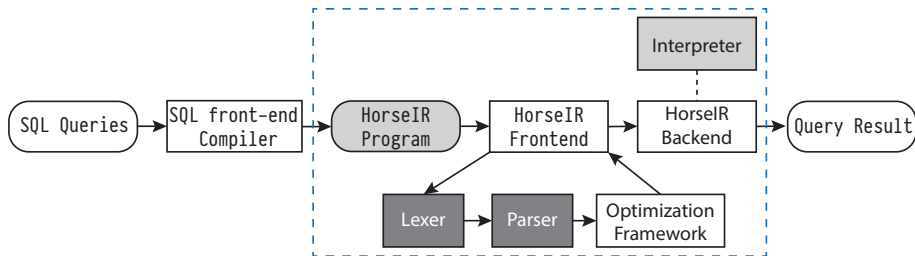


Database System	SQLite	MonetDB	KDB+ System
Language Design	scalar-based	scalar-based	array-based
Storage Structure	row-oriented	column-oriented	column-oriented
Primitives	reduced	complex	reduced
UDF Support	external	MAL, external	native
Parallel Code	limited	explicit declared	implicit (primitives)

# HorseIR - Summary



- ▶ syntax design
- ▶ type system
- ▶ fully functional interpreter connected to Hanfeng's backend





- ▶ array-based IR design for database system
  - ▶ reflects column-data storage in database system
    - ▶ homogeneous data type storage
    - ▶ vector-based column
  - ▶ explore data parallelism easily
    - ▶ a rich set of array-based built-in functions
  - ▶ data compression friendly
- ▶ statically typed, with sub-typing

# HorseIR - A Quick Example



```
SELECT LastName FROM EmployeeTable;
```

```
module default {  
  import Builtin.*;  
  def main() :table {  
    t0    :table      = @load_table('EmployeeTable) ;  
    name  :sym         = check_cast(  
      @column_value(t0 , 'LastName),  
      sym  
    ) ;  
    t1    :list<sym>    = @enlist('LastName) ;  
    t2    :list<sym>    = @enlist(name) ;  
    t3    :table        = @table(t1 , t2) ;  
    return t3 ;  
  }  
}
```



- ▶ column-oriented system (use vector as a abstract view of column)
- ▶ reduced primitive set
  - ▶ efficient implementation
  - ▶ expose details to optimizer
- ▶ UDF implementation:
  - ▶ write UDF in arbitrary language
  - ▶ compile into HorseIR
  - ▶ participate in optimization
- ▶ parallel code generation:
  - ▶ at primitive level
  - ▶ at intra-procedural level
  - ▶ at inter-procedural level



```
CREATE TABLE Employee
(LastName varchar(99), DepartmentID int);
CREATE TABLE Department
(DepartmentID int, DepartmentName varchar(99));

select * from Employee, Department
where Employee.DepartmentID = Department.DepartmentID;
```

# Types in HorselR



- ▶ designed for database systems
  - ▶ string and character
  - ▶ numerics (integer and floating-point)
  - ▶ date and time
  - ▶ functions
  - ▶ table and keyed table
- ▶ parametric polymorphism
  - ▶ list (`list<?>`)
  - ▶ dictionary (`dict<?, ?>`)
  - ▶ enumeration (`enum<?>`)

# Types in HorselR



- ▶ statically checked and inferred at compile time
- ▶ resolve overloading and polymorphic dispatch at compile time
  - ▶ minimize runtime overhead
- ▶ downcast guard elimination

`x : ?`

`castedX : i32 = check_cast(x, i32); (*)`

`avgX : i32 = @Builtin.avg(castedX);`

Is it possible to safely remove this cast? use static analysis



# Types in HorselR - Overloading



- ▶ support more flexible UDF
- ▶ choose the most efficient primitive implementation at compile time
  - ▶ function overloading

```
def foo(x: ?, y :?) :bool {  
  avgX :? = @Builtin.avg(x) ;  
  result :bool = @Builtin.lt(y, avgX);  
  return result;  
}  
@Builtin.avg(?) @Builtin.avg(f64) @Builtin.avg(i32)
```

- ▶ choose which implementation?
  - ▶ in foo(i32, i32),
  - ▶ in foo(f64, i32),
  - ▶ in foo(list<i32>, i32)



- ▶ a lexer and parser
- ▶ interpreter
- ▶ customized abstract syntax tree
- ▶ optimization framework
  - ▶ peephole
  - ▶ static analysis




# Conclusion



## ► HorselR

- array-based IR for database system with powerful primitives
- programming language optimizations for databases
- extensible framework for user-defined functions (UDF)
- efficient parallel code generation



- complete working system
- future experiment on TPC-H benchmarks

# Questions?



```
SELECT *  
FROM T  
WHERE ....
```